

Topic dashboard

```
db <- list(
  title = "Topic dashboard", # heading shown at the top of the page
  stance_rater = "stance_1" # which rater's labels the dashboard shows
)

suppressPackageStartupMessages({
  library(here)
  library(tidyverse) # tidyverse LAST so its verbs win
})
# Used namespaces: jsonlite, glue.
options(knitr.table.max.print = Inf, knitr.kable.max_rows = Inf)
`%||` <- function(x, y) if (is.null(x)) y else x
p_out <- function(f) here::here("outputs", f)
source(here::here("code", "llm_source.R"))
```

Inputs

The dashboard is built only from files already on disk. Each is required except the stance files, which appear one per topic as 03_stance.qmd is run. The codebook is checked against the topic proportions it was drafted from: if 02_topics.qmd has refit since, topics may be renumbered, and a stale codebook would put its labels on the wrong topics. The render stops rather than publish that.

```
needed <- p_out(c("paragraph_theta.csv", "topic_codebook.csv",
                 "topic_summary.csv", "topic_trends.csv"))
invisible(fingerprint(needed)) # stops by name if any input is missing
check_fingerprint(p_out("topic_codebook.csv"), p_out("paragraph_theta.csv"),
  strict = FALSE)

paragraph_theta <- readr::read_csv(needed[1], show_col_types = FALSE)
codebook <- readr::read_csv(needed[2], show_col_types = FALSE) |>
  dplyr::left_join(
    readr::read_csv(needed[3], show_col_types = FALSE) |>
      dplyr::select(topic, frex),
    by = "topic")
trends <- readr::read_csv(needed[4], show_col_types = FALSE)

stance_files <- list.files(here::here("outputs"),
  pattern = "^stance_labels_topic\\d+\\.csv$",
  full.names = TRUE)
stance_all <- stance_files |>
```

```

purrr::map(readr::read_csv, show_col_types = FALSE) |>
purrr::list_rbind()
stance <- if (nrow(stance_all) == 0) NULL else {
  s <- stance_all |> dplyr::filter(rater == db$stance_rater, !is.na(label))
  purrr::set_names(s$label, s$para_uid)
}

writeLines(glue::glue("{nrow(paragraph_theta)} paragraphs, {nrow(codebook)} \\
  topics, trends for {dplyr::n_distinct(trends$topic)} topics, stance labels \\
  from {length(stance_files)} topic file(s)."))

```

1484 paragraphs, 12 topics, trends for 12 topics, stance labels from 1 topic file(s).

Build

```

out <- p_out("topic_dashboard.html")
write_dashboard(
  path = out,
  template = here::here("code", "dashboard_template.html"),
  paragraph_theta = paragraph_theta,
  codebook = codebook,
  trends = trends,
  stance = stance,
  title = db$title)

tibble::tibble(
  item = c("file", "size (MB)", "topics", "paragraphs", "documents",
    "topics with a proposition", "paragraphs with a stance label",
    "built"),
  value = c(basename(out), round(file.size(out) / 1e6, 2), nrow(codebook),
    nrow(paragraph_theta), dplyr::n_distinct(paragraph_theta$atom_id),
    sum(!is.na(codebook$proposition)), length(stance),
    format(Sys.time(), "%Y-%m-%d %H:%M")) |>
tbl_paged(caption = "Dashboard build record.")

```

Dashboard build record.

item	value
file	topic_dashboard.html
size (MB)	0.94
topics	12
paragraphs	1484
documents	60
topics with a proposition	1

item	value
paragraphs with a stance label	104
built	2026-09-24 07:35

Before sharing it

The dashboard is one self-contained file: it opens in any browser, from disk or a web host, with no server and no internet connection. **It embeds the full text of every document**, so publishing it republishes the corpus; confirm that is acceptable for this corpus before posting it anywhere public. The prevalence chart is a model estimate with an interval and the yearly counts are exact, and the page labels them as the different claims they are.